

Reinforcement Learning Report Rubric
CS7641: Machine Learning
Fall 2025

Due Dates

- **11/24/2025** Grading Instructions Email sent to Team.
- **11/26/2025** First Two Calibration Reports on your sheet in Teams graded by EOD before holidays (Wednesday).
- **12/02/2025** Roughly 33-50% of Reports Graded.
- **12/05/2025** Roughly 50-80% of Reports Graded.
- **12/09/2025** Complete with RL Report Grading.
- **End of Term** Grades posted.

Rubric

Total Points: 100 + 5 Extra Credit

As a reminder, please give all or nothing for each point item. If a bullet is worth 3 points, award either 3 or 0 for that bullet. This keeps grading consistent across TAs.

Structured Feedback – CS7641 Fall 2025 RL Report

Continuing this semester, I want to put the point values by the section feedback. I have an example document I would like everyone to use as a template. You will still need to add individualized feedback for each student. Please reach out if you have any questions!

1. Requirements (10 points)

- **7 points total: DOCSTRING-GTusername reproducibility file**
 - **2 points:** Report written in LaTeX on Overleaf, and the project exists.
 - **2 points:** GitHub commit hash from GT Enterprise GitHub. Do not award if the hash/link is missing or non-specific (e.g., repo URL only).
 - **2 points:** Clear instructions to run code (environment, commands, and any package/init notes). If these are only on GitHub and not in the DOCSTRING, remove these 2 points.
 - **1 point:** DOCSTRING is a PDF.
- **1 point:** Report is 8 pages or fewer. Do not read beyond page 8; note to the student if over the limit. :contentReference[oaicite:5]index=5
- **1 point:** Graphs and text are legible at 100% zoom on Canvas. If 3+ figures have unreadable axes/legends at 100%, deduct this point.
- **1 point:** References are consistent (APA, MLA, or IEEE) and used correctly. Deduct if:
 - The style is mixed or misapplied (e.g., free-form URLs in body text instead of references).
 - There is no citation beyond course materials / environment docs. Students must connect to at least one external source. :contentReference[oaicite:6]index=6

2. Hypothesis (8 points)

Initial Hypothesis (4 points)

- **2 points:** States at least one clear, testable hypothesis about algorithm–environment behavior (e.g., “On CartPole, off-policy Q-Learning will reach longer average episode lengths than SARSA under the same ϵ -greedy schedule,” or “PI will require fewer sweeps than VI on Blackjack but be more expensive per sweep.”).
- **2 points:** Grounds the hypothesis in some prior evidence or theory: lecture content, Sutton & Barto, the FAQ, a paper, or well-known RL intuition (discounting, on/off-policy, exploration in stochastic vs near-deterministic domains, etc.). The support must be explicit, not just “this seems reasonable.”

Follow-Through (4 points)

- **2 points:** Returns to the hypothesis in the discussion or conclusion, using the actual experiments to evaluate whether it held or not (even partially).
- **2 points:** Cites specific evidence from the report (plots, tables, metrics) in that follow-up. It is fine if the hypothesis is refuted; what matters is explicit, data-backed reflection rather than a throwaway line.

3. MDP Overview and Discretization (24 points)

- **Blackjack MDP (8 pts)**

- (2 points) State tuple: (player sum, dealer up card, usable-ace flag) correctly described.
- (2 points) Actions: hit / stick.
- (2 points) Reward: +1 win, 0 draw, -1 loss (minor convention variants OK if clearly stated).
- (2 points) Episodic nature explained: termination via bust or sticking, and the notion of an episode clearly articulated.

- **CartPole MDP (8 pts)**

- (2 points) Lists the four continuous state variables (cart position, cart velocity, pole angle, pole angular velocity).
- (2 points) Action space: left / right force.
- (2 points) Reward: +1 per time step until failure.
- (2 points) Termination conditions: angle and position thresholds (e.g., $\pm 12^\circ$ and ± 2.4), or clearly equivalent numeric ranges.

- **Comparison & Discretization Strategy (8 pts)**

- (2 points) Highlights that Blackjack is small/discrete while CartPole is continuous and requires discretization for VI/PI and tabular RL.
- (2 points) Explains the implication: tabular DP is straightforward for Blackjack, while CartPole needs a binned state space.
- (2 points) Describes discretization approach (clamps, number of bins per feature, and how continuous states map to bins).
- (2 points) Justifies bin choices: discusses fidelity vs runtime/memory trade-offs, ideally with order-of-magnitude state counts or simple numerical examples.

4. Value Iteration and Policy Iteration (30 points)

- **Explanation of VI and PI (8 pts)**

- (2 points) VI: describes iterative Bellman optimality updates for all states until convergence.
- (2 points) PI: describes the loop of policy evaluation (compute V^π) and greedy policy improvement.
- (2 points) Notes convergence guarantees for finite MDPs with known transition dynamics.
- (2 points) Clearly contrasts VI vs PI in terms of update style and computational trade-offs (e.g., more sweeps for VI vs heavier iterations for PI).

- **Convergence Visualizations (8 pts)** Four distinct plots: VI and PI on both Blackjack and CartPole.

1. (2 points) VI – Blackjack: ΔV vs iteration (max or mean absolute change in V), showing a monotone or near-monotone decline below a stated threshold.
2. (2 points) PI – Blackjack: number of policy changes per iteration (line or bar plot), with a rapid drop as the policy stabilizes.
3. (2 points) VI – CartPole: ΔV vs iteration on the discretized state space; plot should mention discretization settings or reference them in the caption.
4. (2 points) PI – CartPole: policy-change plot across iterations; ideally notes how discretization coarseness affects stability and convergence.

- **Convergence Analysis (7 pts)**

- (2 points) Hyperparameter variation: explores at least θ and γ (3+ values ideal). If only 2 values for each are shown, award 1 point. If no hyperparameter variation for VI/PI, award 0.
- (2 points) Explains why PI often converges in fewer iterations (due to greedy policy improvement), despite more expensive iterations.
- (1 point) Reports iteration counts and explicit convergence criteria (e.g., $\max_s |V_{k+1}(s) - V_k(s)| < \delta$, or no policy changes).

- (2 points) Integrates quantitative results with theory; if results contradict expectations, discusses plausible causes (e.g., discretization choices, tolerance too loose, numerical issues).

- **Final Policy Comparison (4 pts)**

- (2 points) Clearly states whether VI and PI produced the same policy (for Blackjack at minimum).
- (1 point) If different, gives a concrete reason (e.g., different tie-breaking in arg max, rounding, or looser tolerance). If same, this is essentially a free point.
- (1 point) Refers to a comparative visualization (heatmap, table, or policy map) in the text.

- **State Space Impact on CartPole (3 pts)**

- (1 point) Explicitly acknowledges approximation error due to discretization.
- (2 points) Discusses trade-offs quantitatively or with concrete arguments: more bins \Rightarrow larger state space, potentially better policies, but higher runtime/memory; fewer bins \Rightarrow faster but risk of aliasing and suboptimal policies. Being generous is fine if there's clear reasoning; hardware/time limits are real.

5. Model-Free Control with SARSA and Q-Learning (28 points)

Algorithm Explanation & On-/Off-Policy Nuance (4 pts)

- **4 points (Full):**

- Correctly writes or clearly describes both updates:

$$\text{SARSA} : Q(s, a) \leftarrow Q(s, a) + \alpha [r + \gamma Q(s', a') - Q(s, a)]$$

$$Q\text{-Learning} : Q(s, a) \leftarrow Q(s, a) + \alpha [r + \gamma \max_{a'} Q(s', a') - Q(s, a)]$$

- Explicitly states that SARSA is on-policy (uses the action actually sampled under the current behavior policy) and Q-Learning is off-policy (targets $\max_{a'} Q(s', a')$ regardless of the behavior policy).
- Explains at least one behavioral implication: SARSA tends to be more conservative / risk-averse under ϵ -greedy, while Q-Learning can over-commit to optimistic estimates.

- **2 points (Partial):** Correct equations and “on-policy vs off-policy” labels are present, but behavioral explanation is shallow or fuzzy.

- **1 point (Minimal):** Mentions both algorithms and that they are “different,” but mixes the updates or mislabels on-/off-policy.

- **0:** Missing or fundamentally incorrect.

Learning & Convergence Diagnostics (8 pts)

Goal: show learning and stability for *both* algorithms on *both* environments, not just raw reward curves.

Minimum artifacts:

- At least **four** legible plots:

- SARSA – Blackjack
- Q-Learning – Blackjack
- SARSA – CartPole
- Q-Learning – CartPole

- Each environment must have:

- A performance trajectory (reward or episode length vs episodes).
- Some stability / convergence signal somewhere in the section: e.g., mean $|\Delta Q|$, TD-error decay, rolling variance/std of returns, or an explicit convergence criterion applied and discussed.

Scoring:

- **8 points:** All four (or more) plots present; every algorithm–environment pair has a performance curve; at least two stability-type diagnostics are shown and *interpreted*; each figure has at least one causal explanation (not just “curve goes up”).
- **6 points:** Four curves across the two algorithms and environments, and at least one real stability diagnostic (or an explicit, applied convergence criterion) but coverage or interpretation is thin in one area.
- **4 points:** Curves present but almost entirely descriptive, or one algorithm–environment pair clearly missing.
- **2 points:** ≤ 2 plots, poor labeling, or only raw reward/length curves with no stability evidence and no convergence discussion.
- **0 points:** No meaningful learning plots.

If a student has only performance curves but *clearly states and uses* a convergence criterion (window + threshold) with discussion, you may go up to 6/8 instead of forcing 2/8.

Exploration Strategy Design & Impact (6 pts)

- **Full (6 points):**
 1. Describes at least one exploration scheme (e.g., ϵ -greedy with schedule or Boltzmann) with explicit parameters: initial ϵ , decay horizon/form, floor, temperature, etc. :contentReference[oaicite:11]index=11
 2. Justifies the schedule with respect to both environments (e.g., Blackjack needs persistent exploration of rare states; CartPole can safely reduce exploration once balancing is stable).
 3. Uses quantitative evidence to compare behavior under that schedule (episodes to reach a threshold return, variance after “convergence,” etc.), and applies the same schedule when comparing SARSA vs Q-Learning to isolate on-/off-policy effects.
 4. Connects schedule mechanics to observed behaviors (e.g., “Fast decay on Blackjack caused under-exploration of usable-ace states; SARSA and Q-Learning both plateaued at lower returns.”).
- **4 points:** Strategy and parameters are given, but rationale is superficial (e.g., “we used ϵ -greedy because it is standard”) or evidence is very limited.
- **0 points:** No substantive exploration discussion.

Policy & Behavioral Analysis (6 pts)

- **Full (6 points):**
 1. Compares SARSA vs Q-Learning policies/behaviors on the environments. Uses heatmaps, value tables, or qualitative patterns (e.g., “Q-Learning more aggressively doubles down on high-variance states in Blackjack.”).
 2. Relates these differences to on-/off-policy mechanics, exploration schedule, and variance (not just “they differ”).
 3. Connects at least one RL policy (SARSA or Q-Learning) back to VI/PI policies: matches vs mismatches, with plausible causes (exploration residuals, sampling variance, discretization, tolerance).
- **4 points:** Some comparison is present but the causes are hand-wavy or purely descriptive.
- **2 points:** Differences are listed with essentially no causal reasoning.
- **0 points:** No policy comparison or incorrect explanations.

Robustness & Hyperparameter Sensitivity (4 pts)

- **4 points:**
 - Varies at least three hyperparameters per algorithm (e.g., α , γ , and an exploration parameter such as decay rate or floor), in line with the FAQ requirement that each model validate at least three hyperparameters. :contentReference[oaicite:12]index=12
 - Discusses robustness on *both* environments (qualitative OK for one if compute is tight).
 - States some convergence criterion (e.g., moving-average reward plateau, or mean $|\Delta Q|$ below δ) and uses it in the analysis.
 - Interprets stability vs speed trade-offs (e.g., high α faster initial gains but oscillation; lower γ shorter-sighted behavior in Blackjack).
- **3 points:** Clearly varies at least two hyperparameter families and touches both environments; third is weak or barely addressed.
- **1 point:** Only one hyperparameter explored or analysis is almost entirely qualitative.
- **0 points:** No real sensitivity analysis.

Extra Credit (Up to 5 points) – DQN / Rainbow Component on CartPole

Scoring is discrete: award **5**, **3**, **1**, or **0** only. No half-points.

Context: Extra credit focuses on implementing a DQN-style agent for CartPole (with replay + target network) and, ideally, at least one Rainbow-style component (Double DQN, Dueling, Prioritized Replay, Noisy Nets, n -step, or C51), and comparing it with tabular Q-Learning and SARSA.

5 / 5 (Full Credit) Working DQN with replay buffer and target network; at least one Rainbow component *or* a clearly articulated ablation. Architecture, key hyperparameters (buffer, batch size, learning rate, γ , ϵ schedule, target update) are described. Includes legible learning curves compared to SARSA/Q-Learning, discusses stability issues and mitigation, and provides quantitative comparisons (episodes to threshold, final mean \pm std reward, etc.) with causal explanations.

3 / 5 (Partial) DQN learns on CartPole with at least replay *or* a target network. Architecture/hyperparameters partially described. Some comparison to tabular methods is present but largely qualitative or thin on numbers/mechanisms. Rainbow-style modification may be present but shallowly discussed.

1 / 5 (Minimal / Legacy) Either:

- A serious but unsuccessful DQN attempt (architecture + hyperparameters + description of failure), or
- A tabular Q-Learning extension (e.g., Double Q-Learning) with some comparison to SARSA/DP, but no substantial DQN experiments.

0 / 5 (None) No substantive DQN or Rainbow-related implementation; only vague statements of intent.

Version Control

- 11/24/2025 - TJL Updated assignment to include Q-Learning from summer. Edited the point values to help with homogeny.